

Project BioHealth Clinical Engine (BRD)

Document Type: BRD | Project Code: BIOHEALTH

Domain: Genomic & EHR Data Integration Platform

Lead Contact / Author: Dr. Alistair Finch (Chief Medical Informatics Officer) - a.finch@biohealth-clinical.org

1. Executive Summary & Strategic Objectives

This document represents the formal enterprise specification for **Project BioHealth Clinical Engine**. Developed under the governance of Dr. Alistair Finch, this project addresses mission-critical infrastructure demands within Genomic & EHR Data Integration Platform. The primary mandate of this initiative is to deliver an ultra-resilient, scalable, and secure operational framework.

- BioHealth processes high-throughput genomic sequencing datasets for oncology clinical trials.
- Complies strictly with HIPAA, HITECH, and GDPR data privacy frameworks for patient records.
- HL7 FHIR API interfaces communicate with hospital electronic health record systems nationwide.
- Real-time adverse event detection flags potential drug-drug interaction risks within 300ms.

Attribute	Specification Value
Project Lead Name	Dr. Alistair Finch
Designated Role	Chief Medical Informatics Officer
Official Email	a.finch@biohealth-clinical.org
Classification Level	Enterprise Confidential / Tier 1
Target SLA Uptime	99.999% Availability
Security Standard	ISO/IEC 27001 & SOC2 Type II

2. System Architecture & Technical Specifications

The technical design of Project BioHealth Clinical Engine relies on decoupled microservices and event-driven data distribution. High-frequency telemetry streams are ingested by dedicated edge clusters and routed to central data repositories.

2.1 Subsystem Blueprint

The subsystem pipeline for BIOHEALTH is partitioned into three principal layers:

1. **Ingestion Gateway Subsystem (biohealth-ingest)**: Accepts incoming data payloads over secure TLS/UDP sockets. Maintains strict throughput quotas.
2. **Normalization Subsystem (biohealth-processor)**: Parses incoming data streams, validates schema definitions, and extracts key business attributes.
3. **Persistence Subsystem (biohealth-store)**: Stores normalized events into relational databases and vector indices for real-time RAG querying.

2.2 Data Ingestion Protocol & Hex Mapping

Data packets ingested by Project BioHealth Clinical Engine follow a standardized 256-byte binary schema:

- Bytes 0-3: Magic Header Preamble (0x50524F4A - 'PROJ')
- Bytes 4-7: Subsystem Tenant ID (32-bit Unsigned Integer)
- Bytes 8-15: Precise UTC Timestamp (64-bit Microsecond Epoch)
- Bytes 16-31: Security Authentication Signature (HMAC-SHA256)
- Bytes 32-255: Encrypted Payload Content Array

3. Operational Requirements & Failover Workflows

Ensuring uninterrupted operational availability for Project BioHealth Clinical Engine requires rigorous failover protocols and continuous monitoring. The system incorporates automated health checks executed every 500 milliseconds.

3.1 Disaster Recovery & High Availability

In the event of a primary region disruption, Dr. Alistair Finch has established the following high-availability guidelines:

- **Primary Storage Replication:** Data is synchronously mirrored across three availability zones with zero loss tolerance.
- **Automated Failover Trigger:** If the primary node fails three consecutive health checks, traffic DNS endpoints are automatically reassigned to standby nodes.
- **Emergency Escalation Contact:** In case of severe outage, operational personnel must alert the lead engineer directly.

3.2 Performance Benchmarks

The platform undergoes weekly performance stress testing under simulated load spikes. Key metrics include:

- P99 Processing Latency: Under 25 milliseconds
- Maximum Sustained Concurrency: 100,000 active sessions
- Maximum Memory Overhead: Under 16 Gigabytes per worker pod

4. Stakeholder Directory & Governance

Governance for Project BioHealth Clinical Engine is managed by an interdisciplinary steering committee. Regular monthly reviews ensure compliance with internal risk frameworks and external regulatory requirements.

4.1 Key Personnel Directory

- **Project Sponsor & Lead:** Dr. Alistair Finch
- **Designated Role:** Chief Medical Informatics Officer
- **Direct Communication Channel:** a.finch@biohealth-clinical.org
- **Security & Compliance Lead:** Marcus Vance (m.vance@enterprise-security.org)
- **Operations Steering Coordinator:** Captain Arthur Pendelton (a.pendelton@operations-governance.io)

4.2 Regulatory Compliance Matrix

This project adheres to the following corporate governance standards:

1. **SOC2 Type II Audit Assurance:** Verified controls for security, availability, and confidentiality.
2. **ISO 27001 Certification:** Full compliance with Information Security Management System (ISMS) policies.
3. **GDPR / Data Privacy Mandates:** Right to be forgotten and data minimization routines built into storage schemas.

5. Project Milestones & Implementation Roadmap

The deployment of Project BioHealth Clinical Engine is divided into four distinct execution phases spanning 18 months. Progress is tracked via weekly sprint reviews.

Phase Milestone	Target Quarter	Deliverable Summary	Status
Phase 1: Architecture & PoC	Q1 2026	Core engine prototype & vector index pipeline	Completed
Phase 2: Alpha Testing	Q2 2026	Multi-node benchmark load testing & RBAC audit	In Progress
Phase 3: Beta Field Trial	Q3 2026	Staging deployment with live regional traffic	Scheduled
Phase 4: Full Enterprise Rollout	Q4 2026	Production cutover and legacy system sunset	Planned

5.1 Appendix: Glossary & Contact Escalation

For questions regarding this document or to request modification rights, contact **Dr. Alistair Finch** at a.finch@biohealth-clinical.org. All audit logs for document modifications are permanently recorded in the central database.